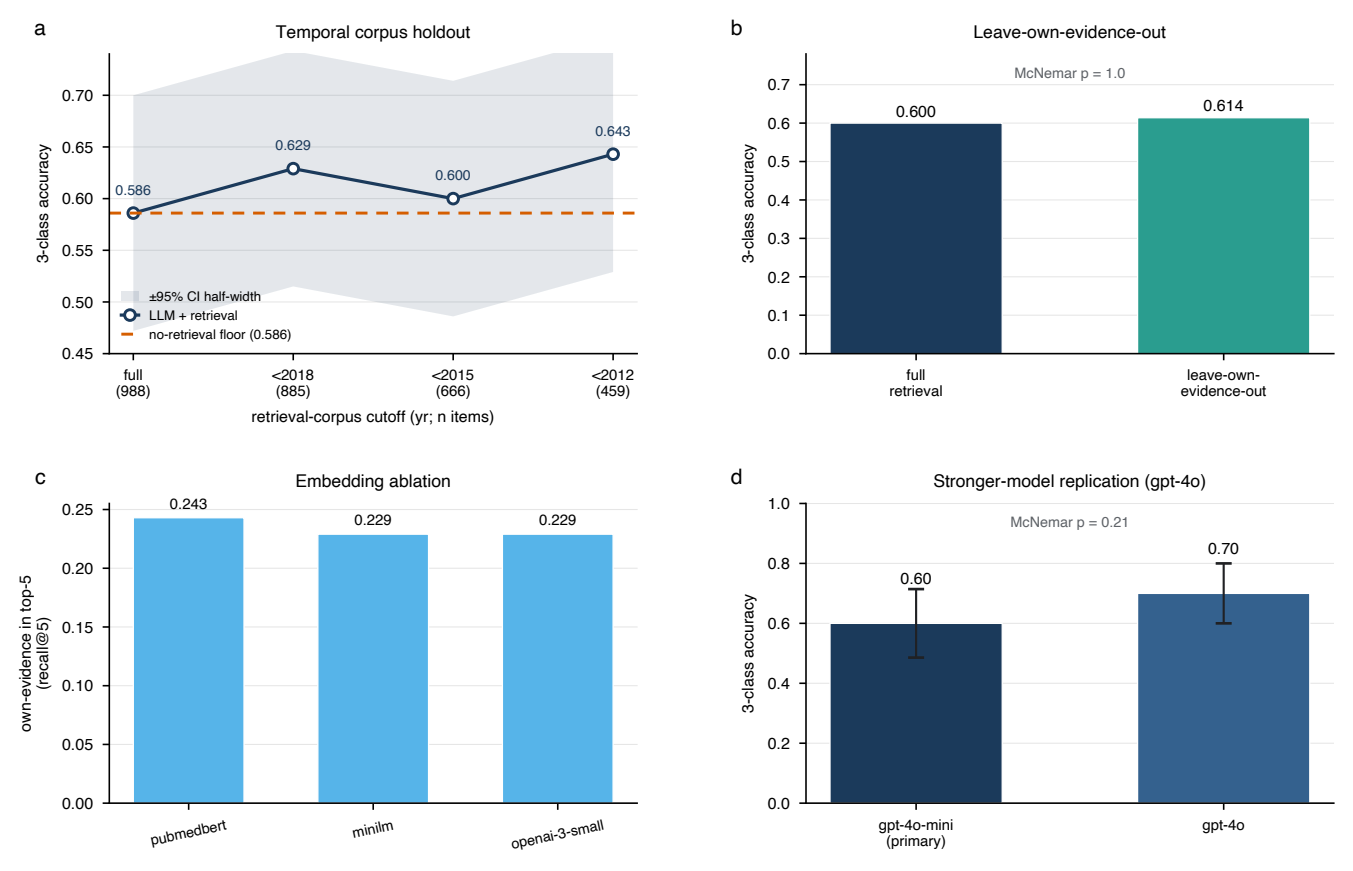
a
Temporal corpus holdout
3-class accuracy
0.70
0.65
0.60
0.55
0.50
0.45
0.586
0.629
0.600
0.643
±95% CI half-width
LLM + retrieval
no-retrieval floor (0.586)
full
(988)
<2018
(885)
<2015
(666)
<2012
(459)
retrieval-corpus cutoff (yr; n items)
b
Leave-own-evidence-out
McNemar p = 1.0
3-class accuracy
0.7
0.6
0.5
0.4
0.3
0.2
0.1
0.0
0.600
0.614
full
retrieval
leave-own-
evidence-out
c
Embedding ablation
own-evidence in top-5
(recall@5)
0.25
0.20
0.15
0.10
0.05
0.00
0.243
0.229
0.229
pubmedbert
minilm
openai-3-small
d
Stronger-model replication (gpt-4o)
McNemar p = 0.21
3-class accuracy
1.0
0.8
0.6
0.4
0.2
0.0
0.60
0.70
gpt-4o-mini
(primary)
gpt-4o